

Parallel Programming

Introduction to CUDA C/C++

Part I

Phạm Trọng Nghĩa

ptnghia@fit.hcmus.edu.vn

Data parallelism

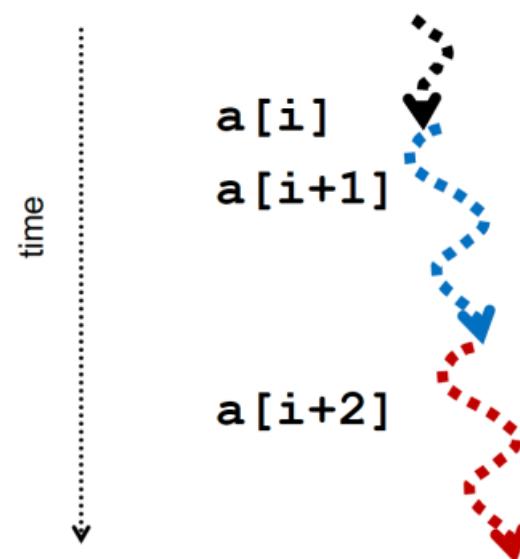
- Question: Why modern software applications run slowly?
- Answer: too much data to process
 - Image-processing apps: million to trillions of pixels
 - Molecular dynamics apps: Thousands to billions of atoms
- Organizing the computation around the data such that we can execute the resulting independent computations in parallel to complete the overall job faster—often much faster.



Sequential Execution Model

- One instruction at the time
- Optimizations possible at the machine level

```
int a[N];
for (i = 0; i < N; i++) {
    a[i] = a[i] * 2;
}
```

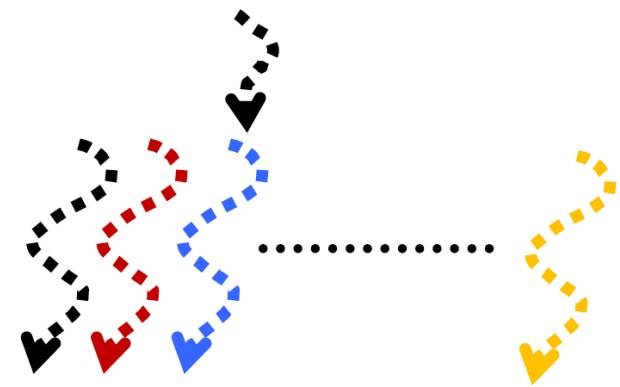


Parallel Execution Model: SIMD

- SIMD: Single Instruction, Multiple Data
- Some instructions executed concurrently

```
int a[N];
for all elements do in parallel{
    a[i] = a[i] * 2;
}
```

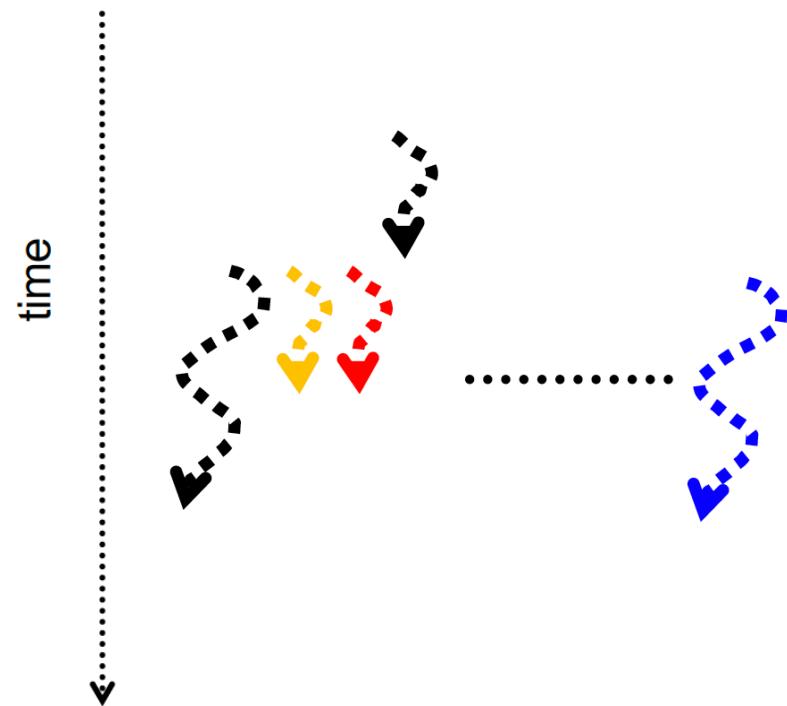
time ↓



Parallel Execution Model: SPMD

- Single program, multiple data
- Code is identical across all threads
- Execution path may differ

```
int a[N];
for all elements do in parallel{
    if (a[i] > threshold)
        a[i] = a[i] * 2;
}
```



CUDA C/C++: is extended-C/C++, allows us to write a program running on both CPU (sequential parts) and GPU (massively parallel parts)

```
#include <iostream>
#include <algorithm>

using namespace std;

#define N 1024
#define RADIUS 3
#define BLOCK_SIZE 16

__global__ void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];
    int gindex = threadIdx.x + blockDim.x * blockDim.x;
    int index = threadIdx.x + RADIUS;

    // Read input elements into shared memory
    temp[index] = in[gindex];
    if (threadIdx.x < RADIUS) {
        temp[index - RADIUS] = in[gindex - RADIUS];
        temp[index + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];
    }

    // Synchronize (ensure all the data is available)
    __syncthreads();

    // Apply the stencil
    int result = 0;
    for (int offset = -RADIUS ; offset <= RADIUS ; offset++)
        result += temp[index + offset];

    // Store the result
    out[gindex] = result;
}

void fill_ints(int *x, int n) {
    fill_n(x, n, 1);
}

int main(void) {
    int *in, *out;           // host copies of a, b, c
    int *d_in, *d_out;       // device copies of a, b, c
    int size = (N + 2*RADIUS) * sizeof(int);

    // Alloc space for host copies and setup values
    in = (int *)malloc(size); fill_ints(in, N + 2*RADIUS);
    out = (int *)malloc(size); fill_ints(out, N + 2*RADIUS);

    // Alloc space for device copies
    cudaMalloc((void **)&d_in, size);
    cudaMalloc((void **)&d_out, size);

    // Copy to device
    cudaMemcpy(d_in, in, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_out, out, size, cudaMemcpyHostToDevice);

    // Launch stencil_1d() kernel on GPU
    stencil_1d<<<N/BLOCK_SIZE,BLOCK_SIZE>>>(d_in + RADIUS,
d_out + RADIUS);

    // Copy result back to host
    cudaMemcpy(out, d_out, size, cudaMemcpyDeviceToHost);

    // Cleanup
    free(in); free(out);
    cudaFree(d_in); cudaFree(d_out);
    return 0;
}
```

parallel function

serial code

parallel code

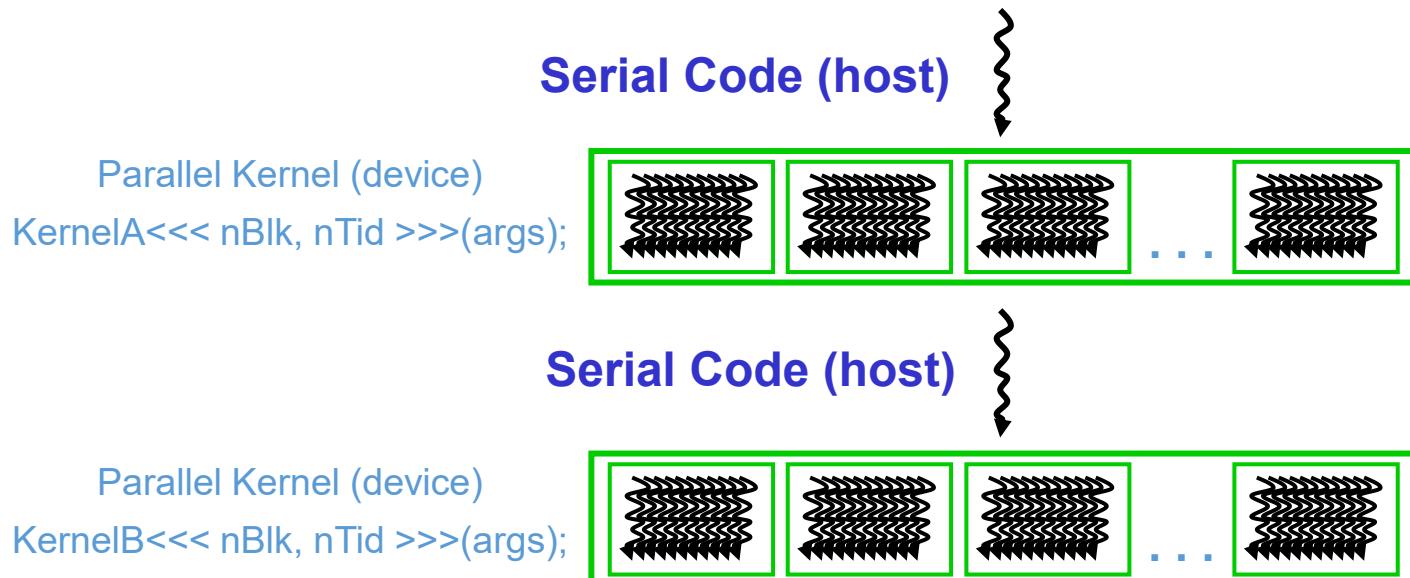
serial code

Host = CPU
(+ memory)

Device = GPU

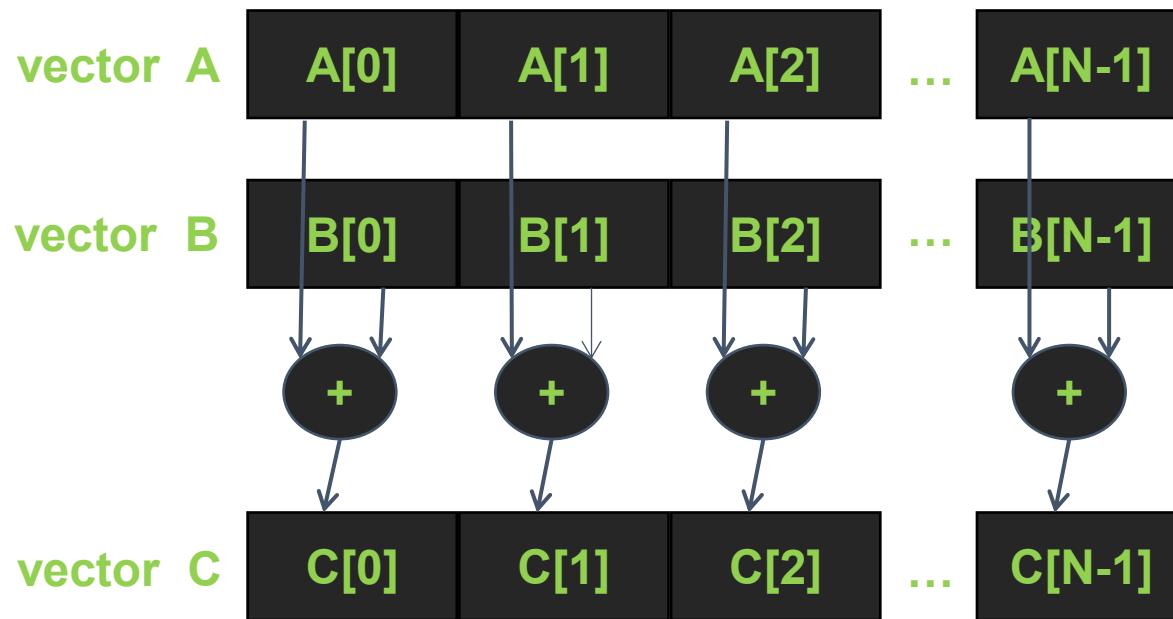
CUDA Execution Model

- Heterogeneous host (CPU) + device (GPU) application C program
 - Serial parts in **host** C code
 - Parallel parts in **device** SPMD kernel code
 - All threads in grid run the same kernel code
 - Each threads has an index that it use to compute memory address and make control decisions



Data Parallelism - Vector Addition

- Adding 2 vectors sequentially using host
- Adding 2 vectors in parallel using device: each thread on device are responsible for computing an element in the sum vector, and all these threads run in parallel



```
int main(int argc, char **argv)
{
    int n; // Vector size
    float *in1, *in2; // Input vectors
    float *out; // Output vector

    // Input data into n
    ...

    // Allocate memories for in1, in2, out
    ...

    // Input data into in1, in2
    ...

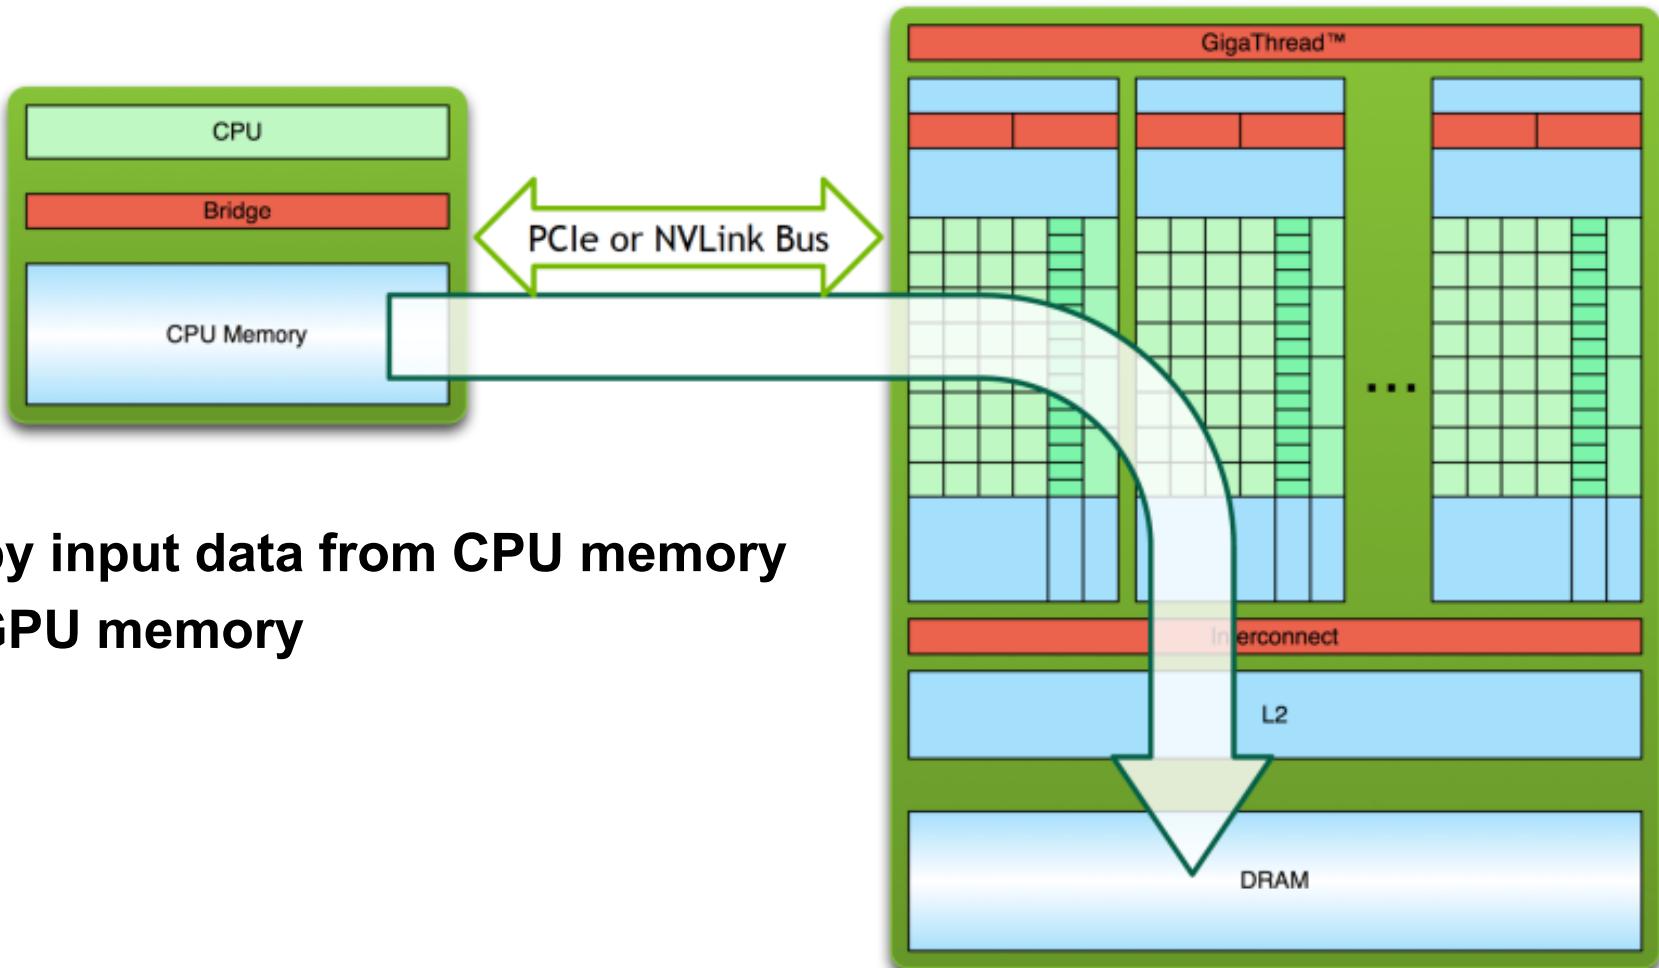
    // Add vectors (on host)
    addVecOnHost(in1, in2, out, n);

    // Free memories
    ...

    return 0;
}
```

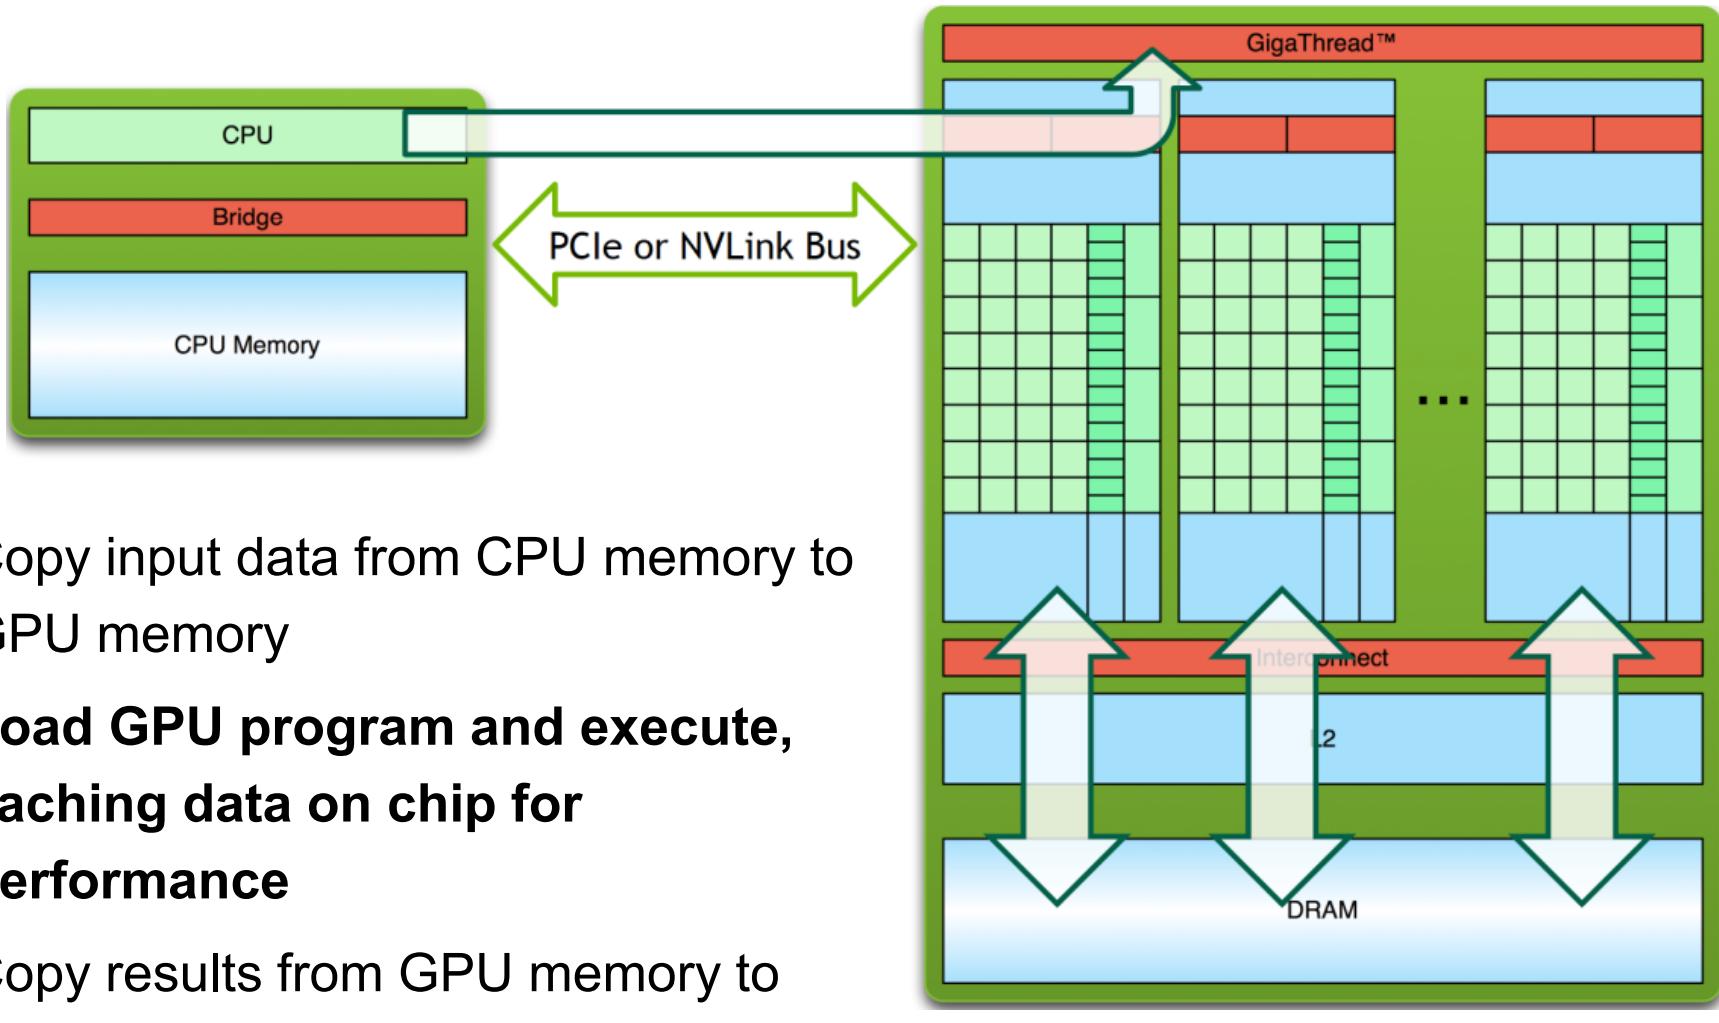
```
void addVecOnHost(float* in1, float* in2, float* out, int n)
{
    for (int i = 0; i < n; i++)
        out[i] = in1[i] + in2[i];
}
```

Simple Processing Flow

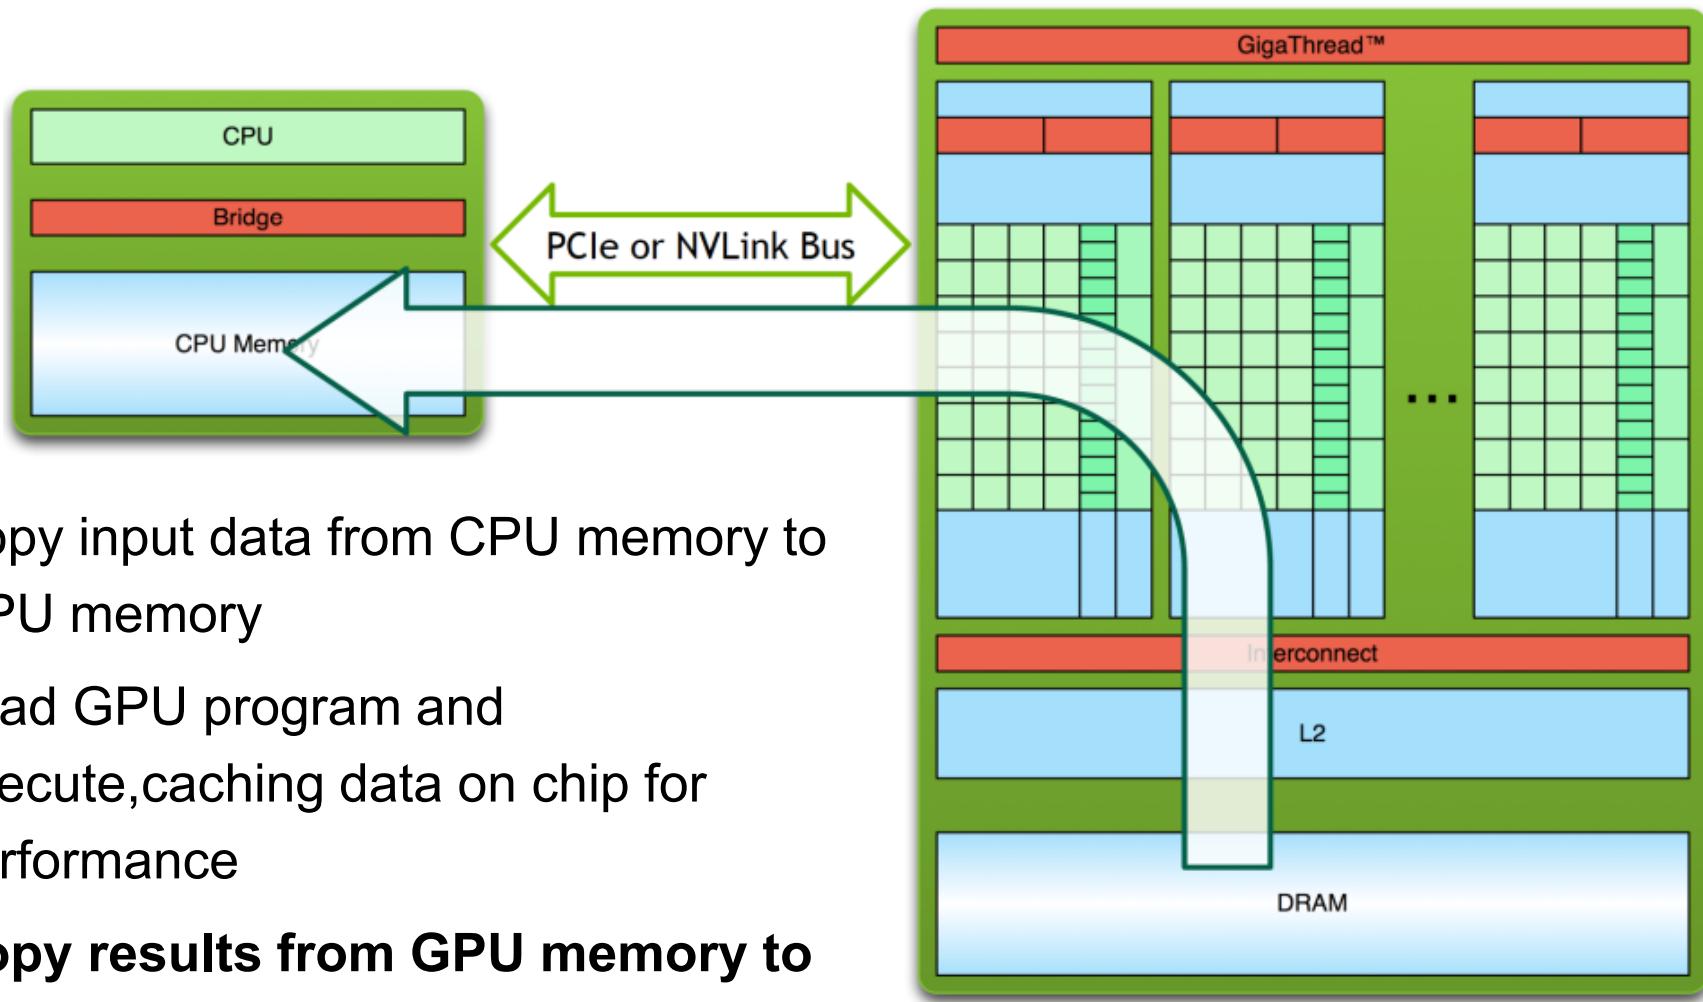


1. Copy input data from CPU memory to GPU memory

Simple Processing Flow



Simple Processing Flow



```

int main(int argc, char **argv)
{
    int n; // Vector size
    float *in1, *in2; // Input vectors
    float *out; // Output vector

    // Input data into n
    ...

    // Allocate memories for in1, in2, out
    ...

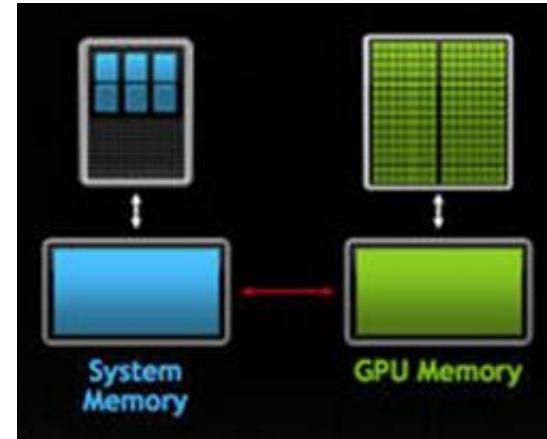
    // Input data into in1, in2
    ...

    // Add vectors (on host)
    addVecOnHost(in1, in2, out, n);

    // Free memories
    ...

    return 0;
}

```



// Host allocates memories on device
...

// Host copies data to device memories
...

// Host invokes kernel function to add vectors on device
...

// Host copies result from device memory
...

// Host frees device memories
...

```
// Host allocates memories on device
float *d_in1, *d_in2, *d_out;
cudaMalloc(&d_in1, n * sizeof(float));
cudaMalloc(&d_in2, n * sizeof(float));
cudaMalloc(&d_out, n * sizeof(float));

// Host copies data to device memories
...
// Host invokes kernel function to add vectors on device
...
// Host copies result from device memory
...
// Host frees device memories
...
```

```
// Host allocates memories on device
float *d_in1, *d_in2, *d_out;
cudaMalloc(&d_in1, n * sizeof(float));
cudaMalloc(&d_in2, n * sizeof(float));
cudaMalloc(&d_out, n * sizeof(float));

// Host copies data to device memories
cudaMemcpy(d_in1, in1, n * sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_in2, in2, n * sizeof(float), cudaMemcpyHostToDevice);

// Host invokes kernel function to add vectors on device
...
// Host copies result from device memory
...
// Host frees device memories
...
```

```
// Host allocates memories on device
float *d_in1, *d_in2, *d_out;
cudaMalloc(&d_in1, n * sizeof(float));
cudaMalloc(&d_in2, n * sizeof(float));
cudaMalloc(&d_out, n * sizeof(float));

// Host copies data to device memories
cudaMemcpy(d_in1, in1, n * sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_in2, in2, n * sizeof(float), cudaMemcpyHostToDevice);

// Host invokes kernel function to add vectors on device
...
// Host copies result from device memory
cudaMemcpy(out, d_out, n * sizeof(float), cudaMemcpyDeviceToHost);

// Host frees device memories
...
```

```
// Host allocates memories on device
float *d_in1, *d_in2, *d_out;
cudaMalloc(&d_in1, n * sizeof(float));
cudaMalloc(&d_in2, n * sizeof(float));
cudaMalloc(&d_out, n * sizeof(float));

// Host copies data to device memories
cudaMemcpy(d_in1, in1, n * sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_in2, in2, n * sizeof(float), cudaMemcpyHostToDevice);

// Host invokes kernel function to add vectors on device
...
// Host copies result from device memory
cudaMemcpy(out, d_out, n * sizeof(float), cudaMemcpyDeviceToHost);

// Host frees device memories
cudaFree(d_in1);
cudaFree(d_in2);
cudaFree(d_out);
```

```
// Host allocates memories on device
float *d_in1, *d_in2, *d_out;
cudaMalloc(&d_in1, n * sizeof(float));
cudaMalloc(&d_in2, n * sizeof(float));
cudaMalloc(&d_out, n * sizeof(float));
```

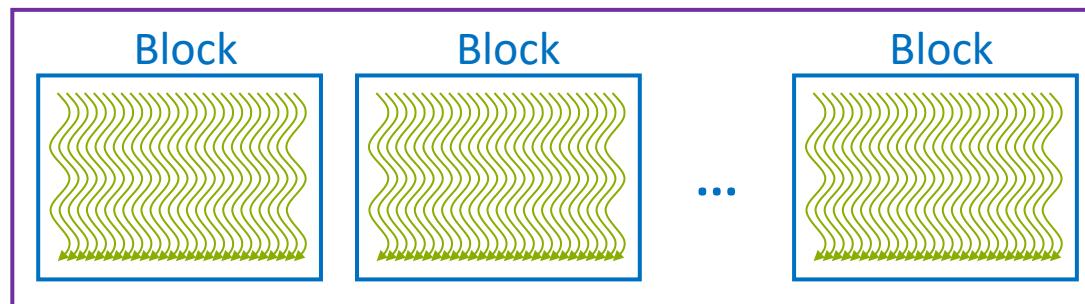
```
// Host copies data to device memories
cudaMemcpy(d_in1, in1, n * sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_in2, in2, n * sizeof(float), cudaMemcpyHostToDevice);
```

// Host invokes kernel function to add vectors on device

```
dim3 blockSize(256); // For simplicity, you can temporarily view blockSize as a number
dim3 gridSize((n - 1) / blockSize.x + 1); // Similarly, view gridSize as a number
addVecOnDevice<<<gridSize, blockSize>>>(d_in1, d_in2, d_out, n);
```

This command creates on device a bunch of threads (called **grid**) executing the addVecOnDevice function in parallel; these threads are organized into gridSize groups or **blocks**, each group/block consists of blockSize threads

Grid



```

...
// Host invokes kernel function to add vectors on device
dim3 blockSize(256);
dim3 gridSize((n - 1) / blockSize.x + 1);
addVecOnDevice<<<gridSize, blockSize>>>(d_in1, d_in2, d_out, n);
...

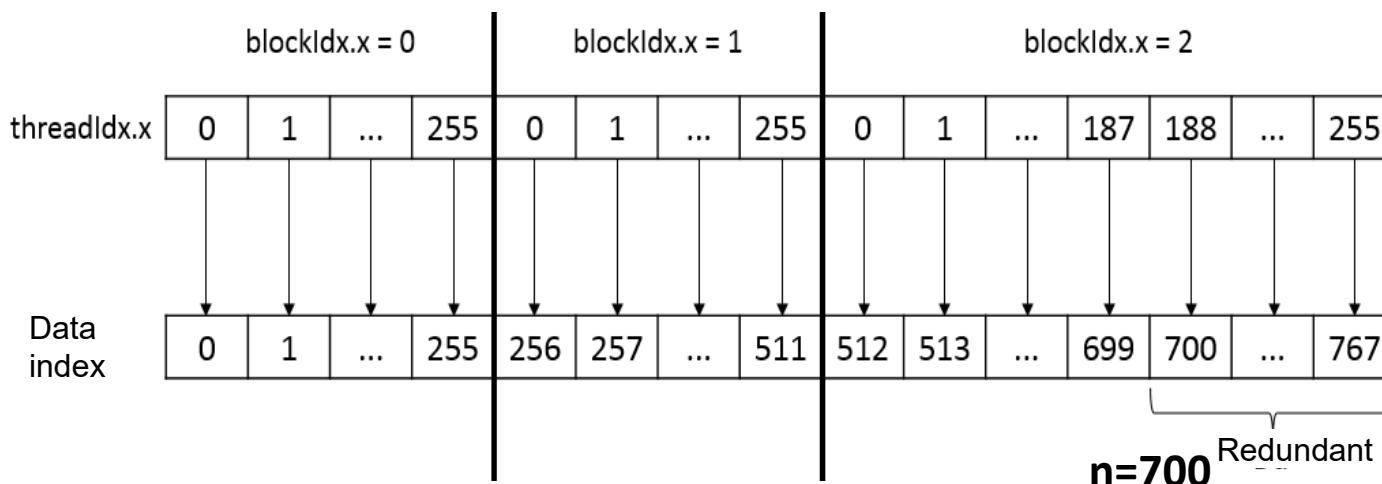
```

Kernel functions must return “void”

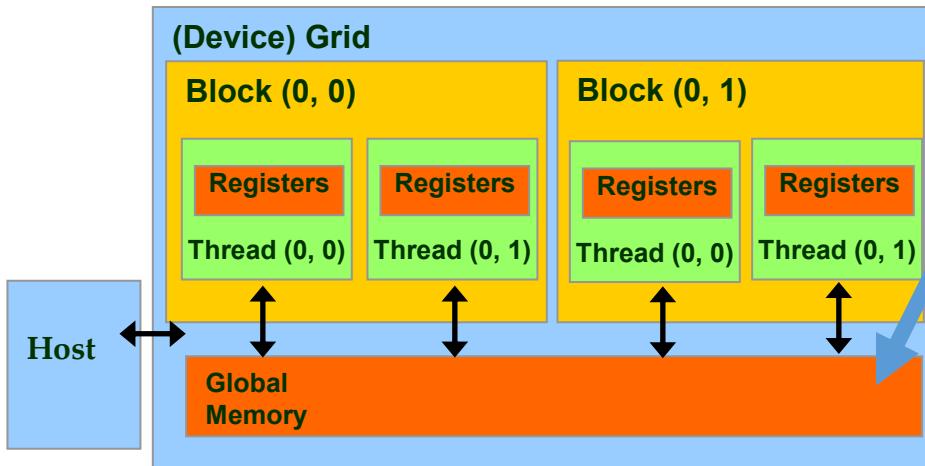
```

__global__ void addVecOnDevice(float* in1, float* in2, float* out, int n)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n)
        out[i] = in1[i] + in2[i];
}

```

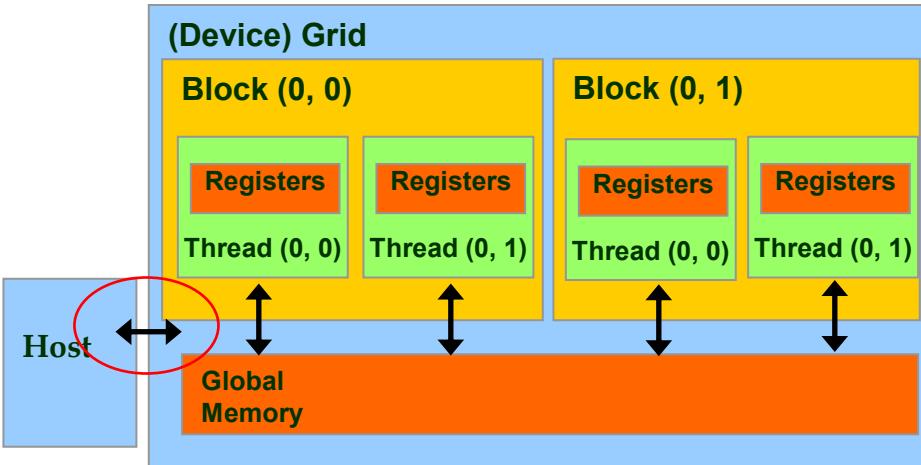


CUDA Device Memory Management API functions



- **cudaMalloc()**
 - Allocates an object in the device global memory
 - Two parameters
 - **Address of a pointer** to the allocated object
 - **Size of** allocated object in terms of bytes
- **cudaFree()**
 - Frees object from device global memory
 - One parameter
 - **Pointer** to freed object

Host-Device Data Transfer API functions



– `cudaMemcpy()`

- memory data transfer
- Requires four parameters
 - Pointer to destination
 - Pointer to source
 - Number of bytes copied
 - **Type/Direction of transfer**
- Transfer to device is **synchronous** with respect to the host

`cudaMemcpyHostToHost` Host -> Host

`cudaMemcpyHostToDevice` Host -> Device

`cudaMemcpyDeviceToHost` Device -> Host

`cudaMemcpyDeviceToDevice` Device -> Device

`cudaMemcpyDefault` Default based unified virtual address space

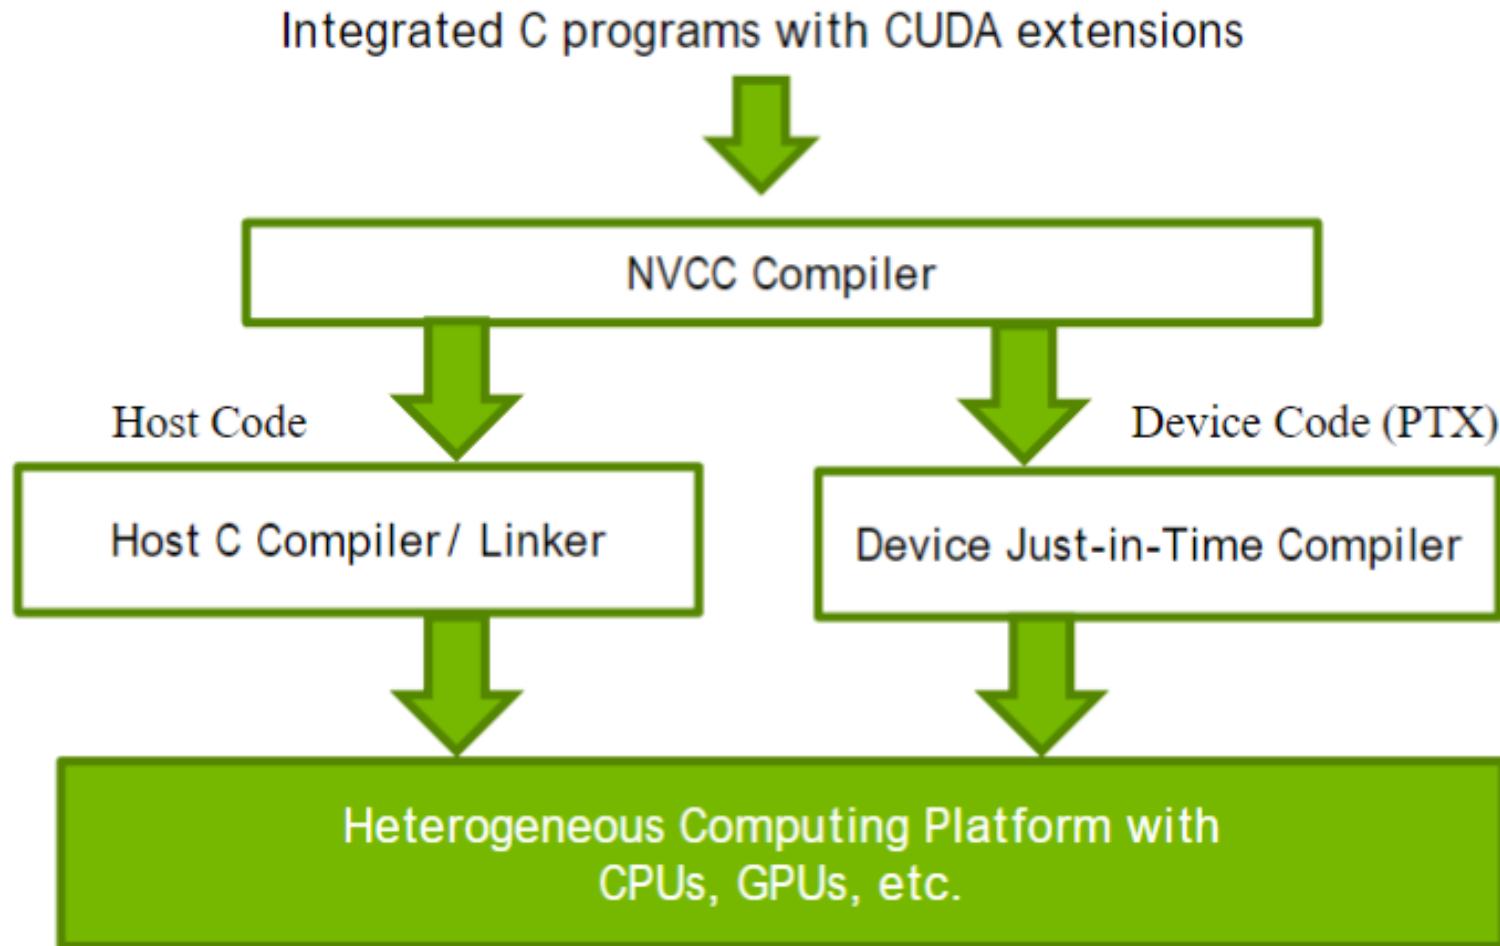
More on CUDA Function Declarations

	Callable from	Execute on	Execute by
<code>__device__ float DeviceFunc()</code>	device	device	Caller host thread
<code>__global__ void KernelFunc()</code>	host	device	New grid of device thread
<code>__host__ float HostFunc()</code>	host	host	Caller thread device

- `__global__` define a kernel function
 - A kernel function must return void
- `__device__` and `__host__` can be used together
 - Generate two versions of object code for the same function
- `__host__` is optional if use alone.

Compiling A CUDA Program

- Use NVCC (NVIDIA C compiler)



Kernel function execution is asynchronous w.r.t host by default

After host calls a kernel function to be executed on device, host will be free to do other works without waiting the kernel to be completed

```
...
// Host invokes kernel function to add vectors on device
dim3 blockSize(256);
dim3 gridSize((n - 1) / blockSize.x + 1);
addVecOnDevice<<<gridSize, blockSize>>>(d_in1, d_in2, d_out, n);

// Host copies result from device memory
cudaMemcpy(out, d_out, n * sizeof(float), cudaMemcpyDeviceToHost); // OK?
```

OK, because the cudaMemcpy function forces host to wait until the kernel finishes, only then it starts to copy

Kernel function execution is asynchronous w.r.t host by default

```
...
// Host invokes kernel function to add vectors on device
dim3 blockSize(256);
dim3 gridSize((n - 1) / blockSize.x + 1);
double start = seconds(); // seconds is my function to get the current time
addVecOnDevice<<<gridSize, blockSize>>>(d_in1, d_in2, d_out, n);
double time = seconds() - start; // OK?
```

...

Kernel function execution is asynchronous w.r.t host by default

```
...
// Host invokes kernel function to add vectors on device
dim3 blockSize(256);
dim3 gridSize((n - 1) / blockSize.x + 1);
double start = seconds(); // seconds is my function to get the current time
addVecOnDevice<<<gridSize, blockSize>>>(d_in1, d_in2, d_out, n);
cudaDeviceSynchronize(); // Host waits here until device completes its work
double time = seconds() - start; // ✓
...
...
```

Error checking when calling CUDA API functions

- It's possible that an error happens but the CUDA program still run normally and give wrong result
 - don't know where to fix bug 😞
 - to know where to fix bug, we should always check error when calling CUDA API functions
- For convenience, we can define a macro to check error and wrap it around CUDA API function calls

```
#define CHECK(call)\\
{\
    const cudaError_t error = call;\\
    if (error != cudaSuccess)\\
    {\\
        fprintf(stderr, "Error: %s:%d, ", __FILE__, __LINE__);\\
        fprintf(stderr, "code: %d, reason: %s\n", error,\\
                cudaGetErrorString(error));\\
        exit(EXIT_FAILURE);\\
    }\\
}
```

```

// Host allocates memories on device
float *d_in1, *d_in2, *d_out;
CHECK(cudaMalloc(&d_in1, n * sizeof(float)));
CHECK(cudaMalloc(&d_in2, n * sizeof(float)));
CHECK(cudaMalloc(&d_out, n * sizeof(float)));

// Host copies data to device memories
CHECK(cudaMemcpy(d_in1, in1, n * sizeof(float), cudaMemcpyHostToDevice));
CHECK(cudaMemcpy(d_in2, in2, n * sizeof(float), cudaMemcpyHostToDevice));

// Host invokes kernel function to add vectors on device
dim3 blockSize(256);
dim3 gridSize((n - 1) / blockSize.x + 1);
addVecOnDevice<<<gridSize, blockSize>>>(d_in1, d_in2, d_out, n);

// Host copies result from device memory
CHECK(cudaMemcpy(out, d_out, n * sizeof(float), cudaMemcpyDeviceToHost));

// Host frees device memories
CHECK(cudaFree(d_in1));
CHECK(cudaFree(d_in2));
CHECK(cudaFree(d_out));

```

Handling kernel errors

- Handling kernel errors is a bit more complicated because kernels execute asynchronously with respect to the host.
- The CUDA runtime maintains an error variable that is overwritten each time an error occurs.
- The function `cudaPeekAtLastError()` returns the value of this variable
- The function `cudaGetLastError()` returns the value of this variable and also resets it to `cudaSuccess`.
- Read [here](#), “Handling CUDA Errors” section, for more info

Handling kernel errors

```
saxpy<<<(N+255)/256, 256>>>(N, 2.0, d_x, d_y);  
cudaError_t errSync = cudaGetLastError();  
cudaError_t errAsync = cudaDeviceSynchronize();  
if (errSync != cudaSuccess)  
    printf("Sync kernel error: %s\n",  
cudaGetErrorString(errSync));  
if (errAsync != cudaSuccess)  
    printf("Async kernel error: %s\n",  
cudaGetErrorString(errAsync));
```

- Synchronous error:
 - Invalid execution configuration parameters, e.g. too many threads per thread block
 - errSync returned by cudaGetLastError().
- Asynchronous errors:
 - On the device after control is returned to the host
 - Out-of-bounds memory accesses

Reference

- [1] Slides from *Illinois-NVIDIA GPU Teaching Kit*
- [2] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj.
Programming Massively Parallel Processors: A Hands-on Approach. Morgan Kaufmann, 2022
- [3] <https://www.eecs.umich.edu/courses/eecs498-APP>



THE END