

Parallel Programming

Introduction to CUDA C/C++

Part II

Phạm Trọng Nghĩa

ptnghia@fit.hcmus.edu.vn

Review: previous lecture

Structure of a simple CUDA program

- Sequential parts will run on host (CPU), massively parallel parts will run on device (GPU)
- From host, to ask device to compute in parallel:
 - Host allocates memory regions on device by calling **cudaMalloc**
 - Host copies necessary data to memory regions on device by calling **cudaMemcpy**
 - Host calls **kernel function**
 - Host copies results from device by calling **cudaMemcpy**
 - Host frees memory regions on device by calling **cudaFree**

Review: previous lesson


Structure of a simple CUDA program

- Sequential parts will run on host (CPU)
- Parallel parts will run on device (GPU)
- From host, to ask device to compute:
 - Host allocates memory regions on device
 - Host copies necessary data to memory regions on device by calling `cudaMemcpy`
 - Host calls **kernel function**
 - Host copies results from device back to host by calling `cudaMemcpy`
 - Host frees memory regions on device

- Kernel function is executed in parallel by many **threads** on device
- These threads are organized as follows:
 - The **grid** consists of **blocks** (these blocks have the same size)
 - Each block consists of **threads**
- When host calls kernel function, host needs to specify grid size and block size
- On device, in kernel function: each thread can use built-in variables (**blockIdx**, **threadIdx**, **blockDim**, **gridDim**) to identify its data portion

Kernel Executes as Grid

- CUDA kernel executes as a grid (array) of threads.
- All threads in grid
 - Run the same kernel code; called a
 - Single Program Multiple Data (SPMD model).
- Each thread has unique index used to
 - Compute memory addresses and
 - Make control decisions.




Today: working with 2D data

- Last lecture:
 - Wrote a kernel function with one-dimensional grid
- This lecture:
 - More generally at *how threads are organized* and learn how threads and blocks can be used to *process multidimensional arrays*
- Multiple examples:
 - Adding 2 matrixes
 - Converting a RGB image to grayscale (intro)
 - Blurring an image (intro)

Multidimensional grid organization


- In general, a grid is a 3D array of blocks, and each block is a 3D array of threads



- Provides a natural way to invoke computation across the elements in a domain such as a *vector*, *matrix*, or *volume*

Kernel Launch Configuration

- Determines how threads are organized and executed in parallel
- Definition of a grid
 - `dim3 gridDim(2, 2, 1);`
 - `dim3 blockDim(4, 2, 2);`
- `KernelFunction<<< gridDim,
blockDim >>>(...);`
 - `gridDim`: the number of blocks in the grid.
 - `blockDim`: the number of threads in each block




Thread Indexing

- Within a kernel, each thread can determine its unique index using ***built-in variables***.
- Build-in variables
 - **threadIdx**: the index of a thread in a block
 - threadIdx.x, threadIdx.y, and threadIdx.z
 - **blockIdx**: the index of a block in the grid
 - blockIdx.x, blockIdx.y, and blockIdx.z
 - **blockDim**: the dimensions of the block
 - blockDim.x, blockDim.y, and blockDim.z
 - **gridDim**: the dimensions of the grid
 - gridDim.x, gridDim.y, and gridDim.z

Thread Indexing

- Build-in variables: **threadIdx, blockIdx, blockDim,girdDim**



- ThreadIdx tuple is unique to each thread **WITHIN A BLOCK**.

Multidimensional grid organization

- Typical host kernel call:

```
dim3 blockSize(16, 16, 2);
dim3 gridSize(128, 1, 1);
addVec<<<gridSize, blockSize >>> (...);
```

```
dim3 blockSize(16);
dim3 gridSize(128);
addVec<<<gridSize, blockSize >>> (...);
```

```
dim3 blockSize(16);
dim3 gridSize(128);
addVec<<<128, 16 >>> (...);
```

- For 2D (and 1D blocks), simply use block dimension 1 for Z (and Y).


Dimension limit

```
Max gridDim.x: 2147483647  
Max gridDim.y: 65535  
Max gridDim.z: 65535  
Max blockDim.x: 1024  
Max blockDim.y: 1024  
Max blockDim.z: 64  
Max thread per block: 1024
```

- (512, 1, 1) ✓
- (8, 16, 4) ✓
- (32, 16, 2) ✓
- (32, 32, 2) ✗

Adding 2 matrixes

- **Sequential implementation:** easy 😊
- **Parallel implementation:** not difficult 😊
 - Let each thread compute one element in the output matrix
 - With 2D data, it's natural to organize threads into 2D blocks and 2D grid



blockIdx.x = 2 threadIdx.x = 3
blockIdx.y = 1 threadIdx.y = 1


$r = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$
= $1 * 16 + 1$
= 17

$c = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
= $2 * 16 + 3$
= 35

Adding 2 matrixes

- **Sequential implementation:** easy ☺
- **Parallel implementation:** not difficult ☺
 - Let each thread compute one element in the output matrix
 - With 2D data, it's natural to organize threads into 2D blocks and 2D grid
- One final consideration before we code: how should we store a matrix?
 - Store as 2D array (a double pointer)
It will cause a lot of troubles (you can try and see ...)
 - Store as 1D array (a single pointer) by concatenating rows together
It will make life easier

Row-major order



When we implement functions operating on 2D array:

- We can **view it directly as 1D array**, but this approach doesn't work in all cases
- A more general approach is to **view it as 2D array**, compute row and column indexes, and when we need to access, convert row and column indexes to indexes in 1D array (we will use this approach)

Adding 2 matrixes: live coding

Today: working with 2D data

- Adding 2 matrixes
- Converting a RGB image to grayscale (intro)
- Blurring an image (intro)

Converting a RGB image to grayscale

Input




Output




Converting a RGB image to grayscale

Input



Output



Each pixel is represented by
3 numbers:

Each pixel is represented by
1 number:


$$\text{red, green, blue} \xrightarrow{\hspace{1cm}} \text{gray}$$
$$\text{gray} = 0.299 * \text{red} + 0.587 * \text{green} + 0.114 * \text{blue}$$

Each number is an unsigned
integer from 0-255

This number is an unsigned
integer from 0-255

Converting a RGB image to grayscale

- The pixels can be calculated independently of each other



Today: working with 2D data


- Adding 2 matrixes
 - Converting a RGB image to grayscale (intro)
 - Blurring an image (intro)
-
- CUDA functions to query device info

Blurring an image

Input



Output



Blurring an image

To blur an RGB image, we blur each channel separately

→ We just need to understand how to blur one channel

How to blur one channel?

Convolve it with a special filter to make adjacent pixels similar (details in the upcoming homework)

Today: working with 2D data

- Adding 2 matrixes
 - Converting a RGB image to grayscale (intro)
 - Blurring an image (intro)
-
- CUDA functions to query device info

Querying device info

In a CUDA program, we may want to query device info:

- For example, we want to query the max number of threads / block and the max number of blocks / grid on device and set block size and grid size appropriate for this device
→ when device is changed, code is unchanged
- Or we want to query info of all available devices in our machine and select the most powerful device to run the program
- Or we simply want to query device info and print to the screen

Querying device info

Query the number of devices

```
int devCount;  
cudaGetDeviceCount(&devCount);
```

**Select a device to use (in case there are many devices),
e.g. device 0**

```
cudaSetDevice(0);
```

Query info of a device, e.g. device 0

```
cudaDeviceProp devProp;  
cudaGetDeviceProperties(&devProp, 0)
```

Check [here](#) for fields of `cudaDeviceProp` struct. E.g.,
`devProp.maxThreadsPerBlock` indicates the max num of threads / block

Colab device's query results (now)

Device 0: Tesla T4

Compute capability: 7.5

GMEM: 15 GB

Max number of threads per block: 1024

Max size of each dimension of a block: 1024 x 1024 x 64

Max size of each dimension of a grid: 2147483647 x 65535 x 65535

...

Reference

- [1] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj.
Programming Massively Parallel Processors: A Hands-on Approach. Morgan Kaufmann, 2022



THE END