

Parallel Programming

Parallel Execution in CUDA

Phạm Trọng Nghĩa

ptnghia@fit.hcmus.edu.vn

Review of previous lecture

CUDA allows us to organize grid and block as 1D or 2D or 3D


Technical Specifications	Compute Capability													
	5	5.2	5	6	6.1	6.2	7	7.2	7.5	8	8.6	8.7	8.9	9
Maximum number of resident grids per device (Concurrent Kernel Execution)	32		16	128	32	16	128	16					128	
Maximum dimensionality of grid of thread blocks									3					
Maximum x -dimension of a grid of thread blocks									$2^{31}-1$					
Maximum y- or z-dimension of a grid of thread blocks									65535					
Maximum dimensionality of thread block									3					
Maximum x- or y-dimensionality of a block									1024					
Maximum z-dimension of a block									64					
Maximum number of threads per block									1024					
Warp size									32					
Maximum number of resident blocks per SM					32				16	32	16	24	32	
Maximum number of resident warps per SM					64				32	64	48		64	
Maximum number of resident threads per SM					2048				1024	2048	1536		2048	
Number of 32-bit registers per SM									64 K					
Maximum number of 32-bit registers per thread block									64 K					
Maximum number of 32-bit registers per thread									255					
Maximum amount of shared memory per SM	64 KB	96 KB	64 KB	96 KB	64 KB	96 KB	64 KB	164 KB	100 KB	164 KB	100 KB	228 KB		
Maximum amount of shared memory per thread block				48 KB		96 KB	96 KB	64 KB	163 KB	99 KB	163 KB	99 KB	227 KB	
Number of shared memory banks							32							
Maximum amount of local memory per thread								512 KB						
Constant memory size								64 KB						
Cache working set per SM for constant memory		8 KB		4 KB					8 KB					2

Today: level 2 CUDA

Aspects of the **GPU compute architecture** that are essential for CUDA C programmers.

- A high-level, simplified view of the **compute architecture** and explore the concepts of flexible *resource assignment*, *scheduling of blocks*, and *occupancy*.
- *Thread scheduling*, *latency tolerance*, *control divergence*, and *synchronization*.
- **API functions** that can be used to query the resources that are available in the GPU and the tools to help estimate the occupancy of the GPU when executing a kernel.


Architecture of a modern GPU



- GPU consists of SMs – Streaming Multiprocessors
 - Each SM consists of SPs – Streaming Processors (or CUDA cores)
 - The SMs have special on-chip memory
- E.g.: the Ampere A100 GPU has 108 SMs with 64 cores each, totaling 6912 cores in the entire GPU

More about SM


- CUDA Cores
- Shared Memory/L1 Cache
- Register File
- Load/Store Units
- Special Function Units
- Warp Scheduler



Fermi SM

Source: Professional CUDA C
programming

Architecture of a modern GPU




Parallel execution – SM level

- CUDA virtualizes GPU hardware architecture
 - Block = virtual SM
 - Thread = virtual SP
- When host calls a kernel function, a grid of blocks will be created and each block (virtual SM) will be assigned to a real SM for execution
 - Each SM can contain > 1 block to execute
 - It depends on SM resource limitations and resources each block needs
 - E.g., SM needs resources (registers) to keep track of indexes of blocks and threads as well as their **execution state**, SM 2.x resources can afford at most 8 blocks and 1536 threads
 - If block size is 512 → SM 2.x can contain **3** blocks

Parallel execution – SM level

- CUDA virtualizes GPU hardware architecture
 - Block = virtual SM
 - Thread = virtual SP
- When host calls a kernel function, a grid of blocks will be created and each block (virtual SM) will be assigned to a real SM for execution
 - Each SM can contain > 1 block to execute
 - Blocks which have not been assigned to SMs will wait in a queue
 - When a block finishes its execution, a block from the queue will be assigned to the available slot in SM



Note: blocks can be assigned to SMs in an arbitrary order

Parallel execution – SM level

Such parallel execution in CUDA:

- Helps achieve **scalability** 😊




Image source: John Cheng et
al. Professional CUDA C
Programming. 2014


Parallel execution – SM level

Such parallel execution in CUDA:

- Helps achieve **scalability** 😊
- But requires blocks to be independent of each other →
threads in different blocks cannot cooperate (synchronize) with each other 😞
 - Assume **thread a** in **block A** wants to use a result from **thread b** in **block B**,
and GPU resources can only afford executing one block at a time
and currently **block A** is being executed
But **block B** only can be executed when **block A** is done 😞
 - Threads in the same block can cooperate with each other by using
`_syncthreads()` CUDA command

Parallel execution – SM level

An example execution of barrier synchronization
(__syncthreads)



Parallel execution – SM level


__syncthreads() it must be executed by all threads in a block

```
__global__ void incorrect_barrier_example(int n){  
    //...  
    if (threadIdx.x % 2 == 0){  
        //...  
        __syncthreads();  
    }  
    else{  
        //...  
        __syncthreads();  
    }  
}
```

- Not all threads in a block are guaranteed to execute either of the barriers → undefined execution behavior


CUDA Thread Blocks

- All threads in a block execute the same kernel program (**SPMD**)
- Programmer declares block:
 - Block size 1 to 1024 concurrent threads
 - Block shape 1D, 2D, or 3D
- Threads within block have **thread index** numbers
- Kernel code uses **thread index** and **block index** to select work and address shared data
- Threads in the same block **share data** and **synchronize** while doing their share of the work
- Threads in different blocks **cannot cooperate**
- Blocks **execute in arbitrary order!**




Parallel execution – Warp

- In SM, with each block, system don't manage and execute each individual thread but **a group of 32 threads** - called **a warp**




Parallel execution – Warp

- One instruction is executed for all threads in a warp (each thread has its own data)
 - This execution model is called **SIMT** (**Single Instruction Multiple Threads**)
- The benefit of this execution model?
 - Help simplify hardware: less resources for control and more resources for arithmetic throughput



All threads in a warp execute the same instruction when selected

Parallel execution – Warp




Parallel execution – Warp

How is a block in SM divided into warps?


- **Block 1D:** 32 consecutive threads form 1 warp (warp 1: threads 0-31, warp 2: threads 32-63, ...)

If block size is **not a multiple of 32** then the last warp will be padded extra threads so that its size is 32, these threads are **useless but still consume resources**

- **Block 2D:** convert to 1D, then divide as 1D



- **Block 3D:** convert to 1D, then divide as 1D



Parallel execution – SM-inside level

- What if threads in a warp cannot execute the same instruction?
- Warp divergence
- Correctness? OK
 - Speed? Hmm...
 - GPUs use **predicated execution**
 - Each thread computes a yes/no answer for each path
 - Multiple paths taken by threads in a warp are executed serially!
 - If all threads in a warp need a barrier synchronization mechanism: **__syncwarp()**


Parallel execution – SM-inside level

Warp divergence example: branching

ALL THREADS EXECUTE BOTH PATHS

(results kept only when predicate is true
for thread)


```
if (threadIdx.x < 24) {  
    A  
}  
else {  
    B  
}  
C
```



Parallel execution – SM-inside level

Warp divergence example: looping

```
...  
for (int i = 0; i <= threadIdx.x; i++)  
{  
    ...  
}  
...
```



Avoiding Branch Divergence

- Make branch granularity a multiple of warp size.

```
__global__ void divergence1(float* A, float* B, float* C, int n){  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i < n){  
        if (i % 2)  
            C[i] = A[i] * B[i];  
        else  
            C[i] = A[i] / B[i];  
    }  
}
```

```
__global__ void divergence2(float* A, float* B, float* C, int n)  
{  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    if (i < n) {  
        if ((i / 32) % 2 == 0)  
            C[i] = A[i] * B[i];  
        else  
            C[i] = A[i] / B[i];  
    }  
}
```

Example: warp divergence

- Task: adding 2 matrixes
 - Matrix size 1000×1000
 - Each thread computes an element in the result matrix
 - Block size 32×32
- How many diverged warps?
 - A. 0
 - B. 1000 ✓
 - C. 1024
 - D. 2000
 - E. I don't know

Example: warp divergence

Task: adding 2 matrixes

- Matrix size 1000×1000
- Each thread computes an element in the result matrix
- Block size 32×32

Execution time of a diverged warp **vs** non-diverged warp?
(we don't consider non-diverged warps in which all 32 threads fail the if condition and do nothing)

- A. Faster
- B. Slower
-  C. Equal
- D. I don't know

Check Divergent with ncu

- Demo!

```
# Check if ncu is available
```

```
!which ncu
```

```
# Profile branch efficiency with ncu
```

```
!ncu --metrics
```

```
smsp_sass_branch_targets_threads_divergent.sum,smsp_sass_branch_targets.sum  
./warp_demo
```


```
# May need to compile without optimize
```

```
!nvcc -G -arch=sm_75 YourCode.cu
```

- Check: *smsp_sass_branch_targets_threads_divergent.sum*


SM and warp

- Each SM contains **warp schedulers**.
- Each warp scheduler has: a register file, 32-bit integer (INT32), 32-bit floating point (FP32), 64-bit floating point (FP64), and tensor core ALUs, load/store units (LD/ST), and a special function unit (SFU).
- Each SM also contains a local **shared memory** as well as a number of texture units that are shared among its warp schedulers.
- Part of the local shared memory is used as an L1 cache (L1\$), the rest can be used to facilitate fast communication between threads that reside on the same SM




Warp scheduling & latency tolerance

- In SM, are warps executed in parallel?
 - Not totally so
 - E.g., Fermi SM (2.x) can contain at most 48 warps (1536 threads), but it has only 32 cores
 - So, in SM how are warps executed exactly?
 - Why does SM contain many warps / threads compared to its execution resources (cores)?




Warp scheduling & latency tolerance

Before continuing, let's review instruction pipeline



Warp scheduling & latency tolerance


- Every clock cycle:
 - Warp scheduler elects eligible warp
 - Schedules instruction from that warp
- **Latency**: #clock cycles between an instruction being issued and being completed
- **Full compute resource utilization** is achieved when *all warp schedulers have an eligible warp at every clock cycle*.
- The latency of each instruction can be **hidden** by issuing other instructions in other resident warps




Warp scheduling & latency tolerance

- Two instruction types:
 - Arithmetic instructions: 10-20 cycles
 - Memory instructions: 400-800 cycles for global memory accesses
- If an instruction has latency of n clock cycles, then scheduler will need $\sim n$ ready instructions (coming from the same warp or other warps) to “hide” latency, keep pipelines full. (**latency tolerance** or **latency hiding**)
 - ➔ desirable for an SM to have many more threads assigned to it than can be simultaneously supported with its execution resources to maximize the chance of finding a warp that is ready to execute at any point in time


Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance




Warp scheduling & latency tolerance



Warp scheduling & latency tolerance



Warp scheduling & latency tolerance



Resource partitioning

- For each SM:
 - The more warp (threads), the better (to hide latency)
 - Maximum: 2048 Threads/SM = 64 warps/SM (*)
 - Not always can reach this maximum due to **resource availability**

$$\text{Occupancy} = \frac{\text{\#warps SM contains}}{\text{\# max warps SM can contain}}$$

(*) Depend on CC



Resource partitioning

- The execution of a warp mainly consists of the following resources:
 - Program counters
 - Registers
 - Shared memory
- The number of blocks and warps that can simultaneously reside on an SM depends on the number of **registers** and amount of **shared memory** available on the SM and required by the kernel.

Resource partitioning

Registers per SM

Kepler: 64K
Fermi: 32K





More threads with fewer registers per thread

Fewer threads with more registers per thread

Shared Memory per SM

Kepler: up to 48K
Fermi: up to 48K



More blocks with less shared memory per block

Fewer blocks with more shared memory per block

Resource partitioning - Example

- A100 GPU: 32 blocks/SM, 64 warps (2048 threads)/SM, 1024 threads/block, and 65,536 registers/SM
- BlockSize = 32 → ? Theads/SM → Occupancy = ?
- BlockSize = 768 → ? Theads/SM → Occupancy = ?
- Kernel uses 64 registers per thread → ? Theads/SM → Occupancy = ?
- 31 registers/thread, 512 threads/block → ? Theads/SM → Occupancy = ?
- 33 registers/thread, 512 threads/block → ? Theads/SM → Occupancy = ?

Resource partitioning - Example

- A100 GPU: 32 blocks/SM, 64 warps (2048 threads)/SM, 1024 threads/block, and 65,536 registers/SM
- BlockSize = 32 → 1024 Theads/SM → Occupancy = 50%
- BlockSize = 768 → 1538 Theads/SM → Occupancy = 75%
- Kernel uses 64 registers per thread → 1024 Theads/SM → Occupancy = 50%
- 31 registers/thread, 512 threads/block → 2048 Theads/SM → Occupancy = 100%
- 33 registers/thread, 512 threads/block → 1538 Theads/SM → Occupancy = 75%

Resource partitioning - Example

- **Tesla T4 GPU**: 16 blocks/SM, 32 warps (1024 threads)/SM, 1024 threads/block, 65,536 registers/SM
- BlockSize = 32 → ? Threads/SM → Occupancy = ?
- BlockSize = 768 → ? Threads/SM → Occupancy = ?
- Kernel uses 64 registers per thread → ? Threads/SM → Occupancy = ?
- 31 registers/thread, 512 threads/block → ? Threads/SM → Occupancy = ?
- 65 registers/thread, 256 threads/block → ? Threads/SM → Occupancy = ?

Resource partitioning - Example

- **Tesla T4 GPU**: 16 blocks/SM, 32 warps (1024 threads)/SM, 1024 threads/block, 65,536 registers/SM
- BlockSize = 32 → 512 Theads/SM → Occupancy = 50%
- BlockSize = 768 → 768 Theads/SM → Occupancy = 75%
- Kernel uses 64 registers per thread → 1024 Theads/SM → Occupancy = 100%
- 31 registers/thread, 512 threads/block → 1024 Theads/SM → Occupancy = 100%
- 65 registers/thread, 256 threads/block → 768 Theads/SM → Occupancy = 75%

Resource partitioning

- Accurate determination of the number of threads running in each SM can be difficult
- Tool: [CUDA Occupancy Calculator](#)

CUDA Occupancy Calculator

Compute Capability version

8.6

CUDA version

11.1

Threads per block

256

Registers per thread

32

Shared memory per block

2048

bytes

Calculate

- *Or code your own Calculator*


Example: # needed warps for Kepler SM

- Each SM has 4 warps schedulers
→ need 4+ warps / SM

In practice, need many more than 4 to hide latency (Kepler SM can contain up to 64 warps)
- With programs whose performance is limited by computation throughput (latency of a computation instruction: 10+ clock cycles)
 - Without ILP (Instruction Level Parallelism – adjacent independent instructions in a warp):
need 4 schedulers × 10+ cycles
= 40+ warps / SM
 - With ILP: may need fewer warps

More about cores

Example: Kepler SM




- **192 fp32 lanes (cores)**
 - fp32 math
 - Simple int32 math (add,min,etc.)
- **64 fp64 lanes**
- **32 SFU lanes**
 - Int32 multiplies, etc.
 - Transcendentals
- **32 LD/ST lanes**
 - GMEM, SMEM, LMEM accesses
- **16 TEX lanes**
 - Texture access
 - Read-only GMEM access

Image source: Paulius Micikevicius. Performance Optimization, GTC2013

More about cores

Example: Kepler SM



- **192 fp32 lanes (cores)**
 - fp32 math
 - Simple int32 math (add,min,etc.)
- **64 fp64 lanes**
- **32 SFU lanes**
 - Int32 multiplies, etc.
 - Transcendentals
- **32 LD/ST lanes**

NVIDIA “core” refers to fp32 core

fp32 cores > # fp64 cores
→ should use 32-bit float when possible

Guide for block size selection

- Block size should be a multiple of 32 (warp size)
- Block size should be selected so that SM has enough warps to hide latency and make use of available resources
 - Occupancy measure: the ratio of # warps SM contains to # max warps SM can contain

Block size should be selected so that occupancy is high

Example: assume SM can contain up to 8 blocks and 1536 threads (48 warps); what block size should we pick: 64, 256, 1024
 - Not necessary: 100% occupancy = max performance
 - Only need enough warp to hide latency and make use of available resources
 - If warps have ILP, we will need fewer warps
 - ...

Guide for block size selection

- High occupancy (75-100%):
 - ✓ Good for memory-bound kernels
 - ✓ Helps hide memory latency
 - X May not help compute-bound kernels
- Lower occupancy (50-75%):
 - ✓ Can be fine if compute-bound
 - ✓ More registers per thread
 - ✓ More shared memory per block

Block size summary

- Always use multiples of 32 (warp size)
- Default starting points:
 - 1D: 256 threads
 - 2D: 16×16 (256 threads)
 - 3D: $8 \times 8 \times 8$ (512 threads)
- Aim for 50-100% occupancy (but profile to verify benefit)
- Memory-bound → higher occupancy (512-1024 threads)
- Compute-bound → moderate occupancy (128-256 threads)
- Check register and shared memory usage
- Profile with real data - theoretical occupancy ≠ performance
- Consider problem size - very small problems may need smaller blocks
- Test multiple sizes - optimal varies by kernel and GPU

Reference

- [1] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022
- [2] Cheng John, Max Grossman, and Ty McKercher. *Professional Cuda C Programming*. John Wiley & Sons, 2014



THE END